

Dear Editor-in-Chief and Editors,

Please find enclosed the manuscript “**Streaming LLM Weights from CXL Memory: Regimes, Placement and Pipeline-Aware Pinning**” for consideration in the *Journal of Parallel and Distributed Computing*. The manuscript was previously submitted, under the title “CAMP: Content-Aware Memory Prefetching for High-Performance CXL-Based Inference”, to the Journal of Systems Architecture (JSA-D-26-01708) and is submitted here through the Elsevier transfer offer; the reviews received there are answered point by point in the attached *Response to Reviewers*, and the manuscript has been rebuilt and substantially rewritten in response.

Problem. Models such as Llama-3-70B (141 GB in FP16) exceed GPU memory. CXL Type-3 memory expanders offer cheap host-attached capacity, but a GPU cannot address them: weights reach HBM through host-initiated DMA at roughly 10–20 GB/s on published prototypes. We ask when streaming weights over this path is viable and what a runtime policy can still gain.

Contributions. (1) A host-mediated transfer model, a closed-form arithmetic-intensity threshold, an implementation-independent lower bound and a regime map: streaming is nearly free only above about 16k tokens per iteration and tens of times slower than resident weights for small-batch decode. (2) An analysis of where policy gains come from: placement of the resident layers matters when compute and transfer are comparable (contiguous sets idle the link), and a tuned FlexGen-style per-layer split is already within a few per cent of our policy on dense uniform models. (3) CAMP, a runtime policy (work-conserving prefetcher with next-use admission and stall-cost-aware pinning) that stays within 1.07 of the lower bound in simulation, is up to $2.8\times$ faster than demand paging, ties the best baseline on dense models and is ahead where unit sizes or reuse are non-uniform (up to $1.97\times$ on encoder–decoder inference). (4) Negative results: the compute-time estimate changes latency by under 1% once calibrated, and LFU is indistinguishable from LRU on cyclic scans. (5) A simulator whose link and overlap logic is checked on a real GPU, with the accuracy of that check reported as it is (mixed: 0.6% in link-bound configurations, up to 25% elsewhere).

Fit with the journal. The work concerns resource management for memory-disaggregated parallel systems and the scheduling of data movement against computation, which falls within the scope of JPDC.

Limitations are stated in the paper. No CXL hardware was measured; the CXL link is simulated with parameters anchored in published measurements and the pipeline logic is validated on a PCIe-attached GPU with pinned host memory standing in for the pool. Mixture-of-Experts models, multi-GPU or multi-tenant pools and CXL switches are not evaluated, and the T4 check leaves part of its own error unexplained. Code, raw results and the hardware micro-benchmark are publicly available (see the Data availability statement).

The manuscript is original, is not under consideration elsewhere, and the author declares no competing interests. The generative-AI declaration required by the journal is included in the manuscript.

Sincerely,

Quang-Vinh Dang
British University Vietnam, Hung Yen, Vietnam
vinh.dq4@buv.edu.vn ORCID 0000-0002-3877-8024